

CptS 451- Introduction to Database Systems

Relational Model and SQL DDL (DMS Ch-3)

Instructor: Sakire Arslan Ay



Outline

- Relational Model
- Constraints on Relations
 - Domain constraints
 - Integrity constraints
 - Key constraints
 - Foreign key constraints
 - Enforcing integrity constraints
- Introduction to SQL

Relational Model

- Where are we?

Database Design: A 6-Step Program

1. Requirements Analysis: what data, apps, critical operations
2. Conceptual DB Design: high-level description of data and constraints – typically using ER model
3. Logical DB Design: conversion into a schema
 - pick a type of DBMS, relational DBMS is most popular and is our focus
4. Schema Refinement: normalization (eliminating redundancy)
5. Physical DB Design: consider workloads, indexes and clustering of data
6. Application/Security Design

ER Model vs. Relational Model

- Both are used to model data
- ER model has many concepts
 - entities, relations, attributes, etc.
 - well-suited for capturing the app. requirements
 - not well-suited for computer implementation (does not even have operations on its structures)
- Relational model
 - just has a single concept: relation (table)
 - world is represented with a collection of tables
 - well-suited for efficient manipulations on computers

Relational Model

- Main idea:
 - Table: **relation**
 - Table header: **relation schema**
 - Column header: **attribute**
 - Row: **tuple**
- Attributes:
 - Each attribute has a **domain**
 - Each attribute is **atomic**.

Employee

ssno	name	salary
111-11-1111	John Yates	80,000
222-22-2222	Vasiliy Bunakov	40,000

Relational Model

Main Concepts:

1. Table: **relation**
2. Column header: **attribute**
3. Row: **tuple**

Relation R:

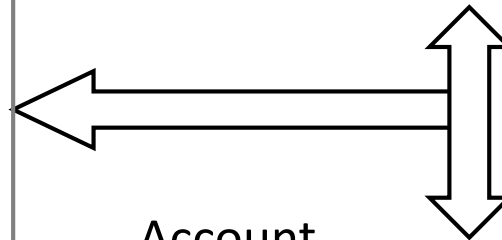
account = { (150, 20, 11000),
(160, 23, 2300),
(180, 23, 32000) }

Relation schema: R(attributes)

Account (AccountId, CustomerId, Balance)

Account

AccountId	CustomerId	Balance
150	20	11,000
160	23	2,300
180	23	32,000



Account

CustomerId	AccountId	Balance
23	160	2,300
23	180	32,000
20	150	11,000

(No order in attributes or tuples)

Relational Model: Some Notation

- The schema (or description) of a relation:
 - Denoted by **R(A1, A2,An)**
 - **R** is the **name of the relation**
 - **A1, A2, ..., An** are the attributes of the relation
 - **R.Ai** refers to **R**'s attribute **Ai**
- A database schema consists of
 - A set of **relation schemas**, e.g., $S = (R1, R2... Rn)$
 - A set of **constraints** over the relation schemas

CourseOfferings		
semester	course#	instructor
Fall, 2013	EE 214	Cole
Fall, 2013	Cpts 121	O'Fallon
Fall, 2013	Cpts422	O'Fallon
Spring,2014	Cpts317	Kalyanaraman
Spring, 2014	Cpts460	Wang

CourseOfferings (semester, course#, instructor)

CourseOfferings.instructor

Relational Model: Some Notation

- $R[A_i]$ denote the column of the relation instance that corresponds to attribute A_i
 - $R[A_1, A_2, A_3]$ refer to the columns for attributes A_1, A_2, A_3
 - If the set of all attributes is $\{A_1, \dots, A_n\}$ then, $R[A_1, \dots, A_n]$ refers to the complete table
- Let t be a tuple (row) in relation R
 - $t[A_i]$ refers to the value of the attribute A_i for tuple t
 - $t[A_1, A_2, A_3]$ refers to the sub-tuple of t containing the values of the attributes A_1, A_2, A_3

CourseOfferings [course#, instructor]
↙

	CourseOfferings		
	semester	course#	instructor
	Fall, 2013	EE 214	Cole
$t \rightarrow$	Fall, 2013	Cpts 121	O'Fallon
	Fall, 2013	CptS422	O'Fallon
	Spring, 2014	CptS317	Kalyanaraman
	Spring, 2014	CptS460	Wang

$t = \langle \text{"Fall, 2013"}, \text{"CptS121"}, \text{"O'Fallon"} \rangle$
 $t[\text{course\#}] = \text{"CptS121"}$
 $t[\text{instructor}] = \text{"O'Fallon"}$

Attributes & Nulls

- Attributes:
 - Each attribute has a **domain**
 - The set of allowed values for each attribute
 - Each attribute is **atomic**
 - We cannot refer to or directly see a subpart of the value
 - Composite values and multi-valued attributes are not allowed
 - The number of attributes is called the **degree** of the relation
- Attributes can have a special value: **NULL**
 - Can mean **not known**: we don't know Jack's address
 - Or, **does not exist**: savings account 1001 does not have "overdraft"

Customer(Id, Name, Addr)

Id	Name	Addr
20	Jared	Pullman
23	Jane	Seattle
32	Jack	NULL

Integrity Constraints over Relations

- Why?
 - Condition specified on a database schema
 - Enable DBMSs to check that new data satisfies the specified conditions
 - Make application programmers' lives easier
 - If DBMS guarantees $\text{account} \geq 0$, the debit application programmer need not worry about overdrawn accounts
 - Enable identification of redundancy in schemas
 - Helps in good database design
 - Help the DBMS in query processing
 - Constraints can help the query optimizer choose a good query execution plan (e.g., by helping it with intermediate result size estimation)

Integrity Constraints over Relations

- Constraint types :
 1. Domain constraints
 2. Key constraints
 3. Entity identity constraints (disallowing null values)
 4. Foreign key constraints
 - Inclusion dependencies (generalization of foreign key constraints)

1. Domain Constraints

- Every attribute has a type (i.e., domain)
 - integer, float, date, boolean, string, etc.
- An attribute can have a **domain constraint**, e.g.:
 - $Id > 0$
 - $Salary > 0.0$
 - $Age < 100$
 - $City \in \{Portland, Seattle, Pullman\}$
- An insertion can try to violate the domain constraint
 - DBMS checks if insertion violates domain constraint and rejects the insertion

Integer	String	String
Id	Name	City
20	Tom	Seattle
23	Jane	San Diego
-2	Jack	Pullman

violations

2. Key Constraints

Keys of Relations:

- Definition:
 - A **minimal** subset of the attributes that uniquely identifies a tuple.
- Logical statement:
 - Given
 - Relation schema: $R = (A_1, A_2, \dots, A_n)$
 - Set of attributes: $\{A_1, A_2, \dots, A_n\}$
 - A subset of attributes: $K \subseteq \{A_1, A_2, \dots, A_n\}$
 - Relational tuples: $T = \{t_1, \dots, t_m\}$
 - For any two distinct tuples $t_1 \in T$ and $t_2 \in T$, if the relation R has the constraint, $t_1[K] \neq t_2[K]$, then K is a **key** of the relation schema R

2. Key Constraints (cont.)

- Example:

Log1(LogID, AccountID, Transact#, Time, Amount)

LogID	AccountID	Transact#	Time	Amount
1001	111	4	1/12/02 12:00:01 PM	\$100
1002	122	4	12/28/01 2:23:30 PM	\$20
1003	333	6	9/1/00 5:03:00 PM	\$60

Log2(AccountID, Transact#, Time, Amount)

AccountID	Transact#	Time	Amount
111	4	1/12/02 12:00:01 PM	\$100
122	4	12/28/01 2:23:30 PM	\$20
333	6	9/1/00 5:03:00 PM	\$60

- Features:

- There are usually multiple *keys*. The set of all keys for a relation is called the *candidate keys*.
- All attributes always form a *key* (Note that the relational model doesn't allow duplicates)
- **Minimal** key (no proper subset is a key)
- If more than one key: choose *one (the minimal)* as the *primary key*

3. Entity Identity Constraints: Disallowing Null Values

- Some fields are too important to contain null values
 - E.g., **Sales(saleID, customer, salesman, date, amount)**
 - Here, we do not want 'customer' to contain a null value
- A primary key **cannot** contain a null value
 - Otherwise, we are not able to differentiate tuples


Sales

saleID	Customer	Salesman	Date	Amount
112587	Tom	Cindy	1/21/2015	2
112588	NULL	Cindy	NULL	3
NULL	NULL	Jack	1/11/2015	4

We don't want these! (Solution: Disallow NULLs for this attribute)

4. Foreign Key Constraints

- Specified between two relations to maintain the correspondence between tuples in these relations

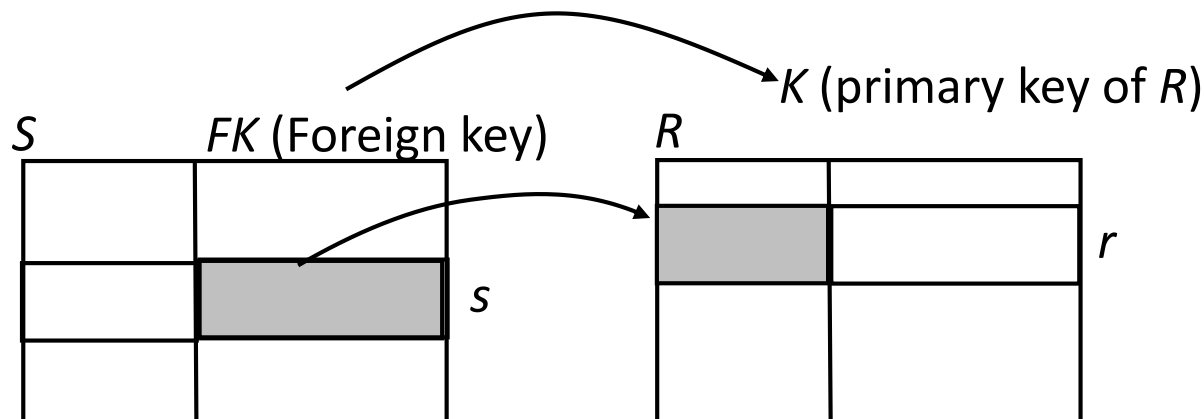


Student.Dept to Dept.Name

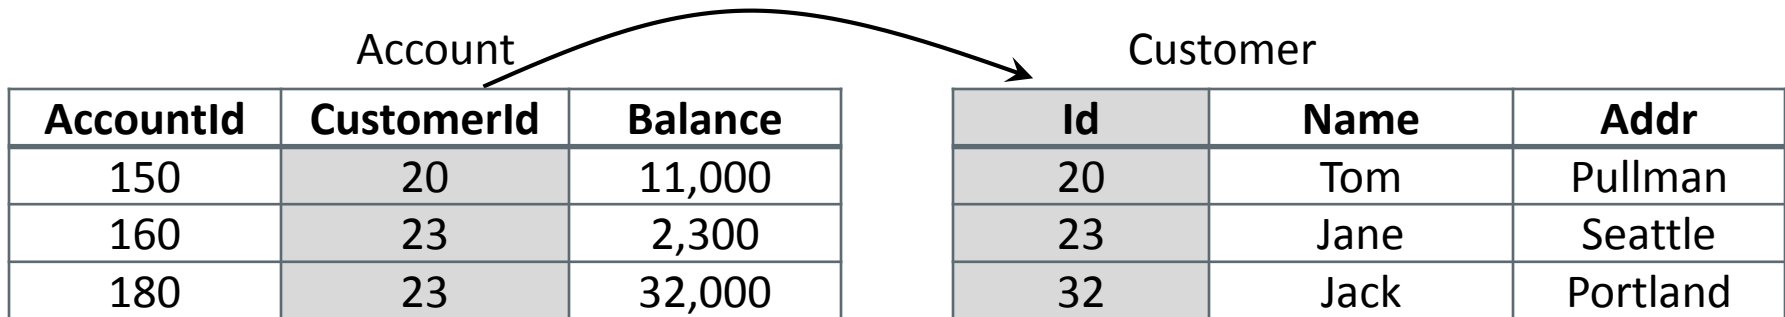
(Every value of ***Student.Dept*** must also be a value of ***Dept.Name***)

Foreign Key Constraints (cont.)

- **Definition** (*foreign key*):
 - Relation schemas: **R** and **S**.
 - **R** has a primary key **K** (a set of attributes)
 - A set of attributes **FK** is a *foreign key* of **S** if:
 - The attributes in **FK** have same domains as the attributes (or correspond to the attributes) in **K**
 - The values for **FK** in each tuple **s** in **S** either occur as values of **K** of a tuple in **r** of **R** or else **FK** is NULL
- **FK** is a foreign key that refers to the primary key **K** of **R**




Another Example



Account.CustomerId to Customer.Id

Inclusion Dependencies

- Generalization of foreign key constraints
 - Inclusion dependency $S[Y_1, \dots, Y_n] \subseteq R[X_1, \dots, X_n]$ means:
values in relation S refer to values in relation R
 - Foreign key is an inclusion dependency in which $\{X_1, \dots, X_n\}$ is the primary key of R



Data Modifications and Constraints

- Modification, insertion, and deletion requests can lead to violations of one or more integrity constraints
- DBMS must check for violations at the end of each operation and suitably allow or disallow the operation
- Impacts of data modifications on foreign key constraints can be sometimes tricky to handle!

Example: Delete from *Referenced* Table

Takes

student#	semester	course#
1111	Fall, 2013	CptS422
2222	Spring, 2014	CptS317
3333	Spring, 2014	CptS317

CourseOfferings

semester	course#	instructor
Fall, 2013	CptS422	O'Fallon
Spring, 2014	CptS317	Kalyanaraman
Fall, 2013	CptS460	Wang

(semester, course#) is the primary key

- Foreign key constraint:
 - Takes(semester, course#) refers to CourseOfferings(semester, course#)
- Deleting a course: “Spring, 2014, CptS317” from “**CourseOfferings**” relation.
- What should happen to any tuples in “**Takes**” that refer to “Spring, 2014, CptS317”?

Integrity Maintenance Policies

1. **Abort**: reject the update (sometimes called **Restrict**)
2. **Cascade**: delete tuples from **Takes** that refer to “Spring,2014 CptS317”
 - Notice that deleting this **Takes** record could have a cascading effect, since other records may refer to it
3. **Set NULL**: set referring values in **Takes** to NULL

student#	semester	course#
1111	Fall, 2013	CptS422
2222	Spring,2014	CptS317
3333	Spring, 2014	CptS317

semester	course#	instructor
Fall, 2013	CptS422	O’Fallon
Spring,2014	CptS317	Kalyanaraman
Fall, 2013	CptS460	Wang

student#	semester	course#
1111	Fall, 2013	CptS422
2222	NULL	NULL
3333	NULL	NULL

semester	course#	instructor
Fall, 2013	CptS422	O’Fallon
Spring,2014	CptS317	Kalyanaraman
Fall, 2013	CptS460	Wang

Example: Update *Referenced* Table



Takes

student#	semester	course#
1111	Fall, 2013	CptS422
2222	Spring, 2014	CptS317 CptS417
3333	Spring, 2014	CptS317 CptS417

CourseOfferings

semester	course#	instructor
Fall, 2013	CptS422	O'Fallon
Spring, 2014	CptS317 CptS417	Kalyanaraman
Fall, 2013	CptS460	Wang

(semester, course#) is the primary key

student#	semester	course#
1111	Fall, 2013	CptS422
2222	NULL	NULL
3333	NULL	NULL

semester	course#	instructor
Fall, 2013	CptS422	O'Fallon
Spring, 2014	CptS417 CptS417	Kalyanaraman
Fall, 2013	CptS460	Wang

- In **CourseOfferings**, change “Spring 2014, CptS317” to “Spring 2014, CptS417”
- Solutions:
 - **Abort (or Restrict)**: reject the update
 - **Cascade**: In **Takes**, also change all appearances of “Spring 2014, CptS317” to “Spring 2014, CptS417”
 - **Set NULL**: set the referencing value(s) to null

Example: Update *Referencing Table*



Takes

student#	semester	course#
1111	Fall, 2013	CptS422 CptS444
2222	Spring, 2014	CptS317
3333	Spring, 2014	CptS317

CourseOfferings

semester	course#	instructor
Fall, 2013	CptS422	O'Fallon
Spring, 2014	CptS317	Kalyanaraman
Fall, 2013	CptS460	Wang

illegal

student#	semester	course#
1111	Fall, 2013	CptS422 CptS460
2222	Spring, 2014	CptS317
3333	Spring, 2014	CptS317

semester	course#	instructor
Fall, 2013	CptS422	O'Fallon
Spring, 2014	CptS317	Kalyanaraman
Fall, 2013	CptS460	Wang

legal

- In **Takes**, change “Fall 2013, CptS422” to a new value
- We need to make sure that the new value belongs to the values of **CourseOfferings**
- If it doesn't, the update should be rejected

Example: Insert to *Referencing Table*



Takes

student#	semester	course#
1111	Fall, 2013	CptS422
2222	Spring, 2014	CptS317
3333	Spring, 2014	CptS317
4444	Fall, 2014	CptS422

CourseOfferings

semester	course#	instructor
Fall, 2013	CptS422	O'Fallon
Spring, 2014	CptS317	Kalyanaraman
Fall, 2013	CptS460	Wang

- In **Takes**, insert tuple “**4444, Fall 2013, CptS422**”
- We need to make sure that the new value for (semester, course#) appears in **CourseOfferings**
- If it doesn't, the insert should be rejected

Introduction to SQL

SQL - Historical Perspective

- Developed by IBM in mid 70s
- First standardized in 1986 by ANSI --- SQL1
- Revised in 1992 --- SQL2. Approximate 580 pages describing syntax and semantics
- In 1999, ANSI/ISO released the SQL3. Many additions for:
 - support for multimedia data
 - addition of abstract data types and object-orientation
 - support for calling programmed functions from within SQL
- Every vendor has a slightly different version of SQL
- If you ignore the details, basic SQL is very simple and declarative. Hence easy to use.

Why SQL?

- SQL is a declarative language.
 - Say “what to do” rather than “how to do it.”
 - Avoid a lot of data-manipulation details needed in procedural languages like C++ or Java.
- Database management system figures out “best” way to execute query.
 - Called “query optimization.”

SQL in Different Roles

- Data definition Language (DDL):
 - allows users to create the relations and constraints.
- Data Manipulation Language (DML)
 - Query Language:
 - declarative -- you specify what you want and not how to retrieve, easy to learn and program
 - Updates:
 - insert, delete, and update tables
- Other commands:
 - View definition language:
 - commands to define rules
 - updates through views generally not supported
 - Transaction Control:
 - commands to specify beginning and end of transactions.

SQL (cont)

- SQL statements can be run interactively via GUI or prompt.
- Embedded SQL:
 - has been embedded in a variety of host languages:
 - C, C++, PERL, Smalltalk, Java (vendor dependent)

Data in SQL

1. Atomic types, a.k.a. data types
2. Tables built from atomic types
3. No nested tables, only flat tables are allowed!
 - We will see later how to decompose complex structures into multiple flat tables

Data Types in SQL

- Characters:
 - CHAR(20) -- fixed length
 - VARCHAR(40) -- variable length
- Numbers:
 - BIGINT, INT, SMALLINT, TINYINT
 - REAL, FLOAT -- differ in precision
 - MONEY
- Times and dates:
 - DATE
 - TIME
 - DATETIME -- SQL Server
- Others...

SQL as DDL (Data Definition Language)

- Simple Data Declarations

```
CREATE TABLE Dept
(dno      int NOT NULL,
dname     varchar(30),
mgr       varchar(20)
);
```

Dept

dno	Dname	Mgr
2752	EECS	Behrooz Shirazi
2534	MME	Michael Kessler
2651	MATH	NULL

Don't allow null values

```
CREATE TABLE Emp
(ename     char(20) NOT NULL,
dno        int DEFAULT 0,
sal        int
);
```

Emp

ename	dno	sal
John Yates	2752	81,000
Michael Kessler	2534	100,500
John Lee	0	32,000

Default value is 0

Define New Domains

```
CREATE DOMAIN personDom CHAR(20);
```

```
CREATE TABLE emp  
( ename      personDom,  
  dno        int default 0,  
  sal        real  
);
```

Creating and Modifying Relations using SQL



```
CREATE TABLE R (  
    A1 D1 [not null] [default V1] [primary key] [unique],  
    ...  
    An Dn [not null] [default Vn] [primary key] [unique],  
    <integrity constraint 1> ,  
    ...  
    <integrity constraint k>  
)
```

Each integrity constraints could be:

1. primary key
2. UNIQUE
3. “check(predicate)”: specifies condition to be satisfied by each tuple
4. foreign key

1. Declaring Keys – Primary Key

There are 2 ways to declare the primary key:

1) When the attribute is listed in the relation schema.

```
CREATE TABLE emp
( ssn      int Primary Key,
  name     char(15),
  dno      int,
);
```

2) As an additional declaration that lists the attribute(s) that form the key.

```
CREATE TABLE emp
( ssn      int,
  name     char(15),
  dno      int,
  Primary Key (ssn)
);
```

```
CREATE TABLE emp
( ssn      int,
  name     char(15),
  dno      int,
  Primary Key (dno,name)
);
```

Need to use
the second
method if the
key consists
of multiple
attributes

2. UNIQUE Constraint

- Alternative key declaration with **UNIQUE**

```
CREATE TABLE emp
( SSN      int Primary Key,
  eID      int UNIQUE,
  name     char(15),
  dno      int,
);
```

```
CREATE TABLE emp
( SSN      int Primary Key,
  name     char(15),
  dno      int,
  UNIQUE (dno,name)
);
```

- A table has only one “primary key” in a table, but many “unique”s
- “Unique” may have NULLs.

3. CHECK Clause

- Enforces a predicate (condition)

```
CREATE TABLE emp  
( SSN      int,  
  name     varchar(25),  
  dno      int,  
  salary   bigint,  
  CHECK (salary >= 0)  
);
```

Ensure that the values of the
assets are non-negative

- Can be named (useful to indicate which constraint an update violated)

```
CONSTRAINT salary_nonzero CHECK (salary >= 0)
```

4. Foreign-Key Constraint

There are 2 ways to declare the foreign key:

1)

```
CREATE TABLE emp
( SSN int,
  name char(15),
  dno int REFERENCES dept(dept#)
);
```

2)

```
CREATE TABLE emp
( ssn int,
  name char(15),
  dno int,
  FOREIGN KEY (dno) REFERENCES dept(dept#)
);
```

- Allow multiple attributes in one foreign-key constraint.
- Allow multiple foreign-key constraints in one table.

emp (SSN, name, dno)

SSN	name	dno
9999	Jack	111
8888	Alice	111
7777	Lisa	222
6666	Tom	333
5555	John	333

dept(dept#, dname, mgr)

dept#	dname	mgr
111	Marketing	9999
222	HR	7777
7777	IT	5555

```
CREATE TABLE dept
( dept# int PRIMARY KEY,
  dname char(15),
  mgr char(15) REFERENCES emp(SSN)
)
```

4. Foreign-Key Constraints (cont.)

- Also called “referential integrity”
- Within an attribute:
 - REFERENCES <TABLE> (<attributes>)
- Separate declaration:
 - FOREIGN KEY <attributes> REFERENCES <table> (<attributes>)
- NULL values allowed for attributes in a foreign key.
 - A foreign-key constraint is automatically satisfied if even one attribute in the foreign key is null.

Giving Names to Constraints

- Keyword **CONSTRAINT**
 - Name a constraint
 - Useful for later altering

```
CONSTRAINT FK1 FOREIGN KEY (dno) REFERENCES dept(dept#);
```

Enforcing Foreign-Key Constraints

emp (SSN, name, dno)

SSN	name	dno
9999	Jack	111
8888	Alice	111
7777	Lisa	222
6666	Tom	333
5555	John	333 444

dept(dept#, dname, mgr)

dept#	dname	mgr
111	Marketing	9999
222	HR	7777
333	Operations	6666
777	IT	8888

emp.dno references dept.dept#

- Modification (insert, update) of **emp**
 - If the new tuple's dno does not exist in dept.dept# , then REJECT!

Enforcing Foreign-Key Constraints

emp (SSN, name, dno)

SSN	name	dno
9999	Jack	111
8888	Alice	111
7777	Lisa	222
6666	Tom	333
5555	John	333

dept(dept#, dname, mgr)

dept#	dname	mgr
111	Marketing	9999
222	HR	7777
333	Operations	6666
777	IT	8888

emp.dno references dept.dept#

- Modification (delete, update) of **dept** whose old “dept#” is referenced by a record in *emp*. There are 3 strategies:
 - Default: reject
 - Cascade: change the *emp* tuple correspondingly
 - Delete in *dept*: delete the referring record(s) in *emp*
 - Update in *dept*: update the dno of the referring record(s) in *emp* to the new dno
 - Set Null: change dno value in referring record(s) in *emp* to NULL

Foreign Key Constraint – Choosing a Policy

- Add “**ON [DELETE,UPDATE] [CASCADE, SET NULL, NO ACTION]**” when declaring a foreign key
 - If no policy is specified, the default (**NO ACTION**) is used.
- Which policy to choose depends on the application.

```
CREATE TABLE emp
( ssn INT,
  name VARCHAR(20),
  dno INT,
  FOREIGN KEY dno REFERENCES dept(dept#)
  ON DELETE SET NULL
  ON UPDATE CASCADE
);
```

SQL as DDL - DROP TABLE command

- Used to remove a relation & its *definition*
 - The relation can no longer be used in queries, updates, or any other commands since its description no longer exists

```
DROP TABLE emp;
```

SQL as DDL - ALTER TABLE command

- Used to add/drop attributes from a relation
- Add attribute:

ALTER TABLE relationName **ADD** attribName attribDomain

- All tuples in the relation are assigned *null as the default value of the new attribute*

- Drop attribute:

ALTER TABLE relationName **DROP** attribName

- Dropping of attributes not supported by many DBMSs

SQL as DDL - ALTER TABLE command (cont.)



- Since new attribute will have NULL values right after the ALTER command is executed, the NOT NULL constraint is not allowed for such an attribute
- Example:

ALTER TABLE emp ADD salary bigint

- The database users must still enter a value for the new attribute “salary” for each EMPLOYEE tuple. This can be done using the UPDATE command. (We will cover this later)