

CptS 451- Introduction to Database Systems

SQL as a Query Language – part2 (DMS ch-5)

Instructor: Sakire Arslan Ay



Nested Queries

- Also called **subqueries**. Embedded inside an **outer** query.
- Similar to function calls in programming languages.
- **Example:** Who is in Sally's department?

```
SELECT E1.ename
FROM   Emp E1, Emp E2
WHERE  E2.ename = 'Sally' AND E1.dno=E2.dno;
```

OR:

```
SELECT ename
FROM   Emp E1
WHERE  Emp.dno IN
      ( SELECT dno
        FROM Emp   E2
        WHERE ename = 'Sally');
```

} subquery

names are scoped

- Semantics:
 - The nested query returns a relation containing dno for which Sally works
 - For each tuple in Emp, evaluate the nested query and check if Emp.dno appears in the set of dno's returned by nested query.

Subqueries - Conditions Involving Relations



- Usually nested queries produce a relation as an answer.
- Conditions involving relations :

1.	s IN R	s is equal to <i>one</i> of the values in R (s is a list of attributes, R is a relation)
2.	s > ALL R	s is greater than <i>every</i> value in unary relation R (s is a single attribute, R is a relation)
3.	s > ANY R, s > SOME R	s is greater than <i>at least 1</i> element in unary relation R. <i>-any</i> and <i>some</i> are synonyms in SQL (s is a single attribute, R is a relation)
4.	EXISTS R	R is <i>not empty</i> (R is a relation)
	Other operators (<, =, <=, >=, <>) could be used instead of >	
	EXISTS can be <i>negated</i> .	

Example 1

- Find the employees with the highest salary.

```
SELECT ename  
FROM emp  
WHERE sal >= ALL  
      (SELECT MAX(sal)  
       FROM Emp);
```

- < ALL, <= ALL, >= ALL, = ALL, <> ALL also permitted

Example 2

- Who makes more than someone in the R&D department?

```
SELECT ename
FROM Emp
WHERE sal > ANY
      (SELECT sal
       FROM Emp, Dept
       WHERE Emp.dno = Dept.dno AND Dept.dname = 'R&D');
```

- “< ANY, <= ANY, >= ANY, > ANY = ANY, <> ANY” are permitted

Nested Queries Producing One Value

- Sometimes subqueries produce a single value

```
SELECT  ename  
FROM    Emp  
WHERE   Emp.dno =  
        (SELECT dno  
         FROM   Dept  
         WHERE  (dname='Production'));
```

- Assume there is only one department called “Production,” then the subquery returns one value.
- If it returns more, it’s a run-time error.

Subqueries-Testing Empty Relations

- “**EXISTS**” checks for nonempty set
- Similarly, “NOT EXISTS” can be used.
- “Find employees who make more money than *some* manager”

```
SELECT ename
FROM   Emp E1
WHERE  EXISTS
      (SELECT ename
       FROM Emp, Dept
       WHERE (Emp.ename = Dept.mgr)
            AND (E1.sal > Emp.sal));
```

E1(ename,dno,salary)

ename	dno	sal
Jack	111	81K
Alice	111	70K
Lisa	222	32K
Tom	333	56K
Mary	333	65K

Emp(ename,dno,salary)

ename	dno	sal
Jack	111	81K
Alice	111	70K
Lisa	222	32K
Jack	333	81K
Tom	333	56K
Mary	333	65K

Dept(dno,dname, mgr)

dno	dname	mgr
111	Sells	Alice
222	Toys	Lisa
333	Electronics	Mary

Result

E1.ename
Jack
Alice
Tom
Mary

Subqueries-Testing Empty Relations (cont)



- The nested query uses attributes name of E1 defined in outer query. These two queries are called *correlated*.
 - Semantics: for each assignment of a value to some term in the subquery that comes from a tuple variable outside, the subquery needs to be executed

Subqueries-Testing Empty Relations

```
SELECT ename
FROM Emp as E1
WHERE EXISTS
    (SELECT ename
     FROM Emp, Dept
     WHERE (Emp.ename = Dept.mgr)
           AND (E1.sal > Emp.sal));
```

- Evaluate **exists** condition on (Jack,111,81K)
 - Subquery returns: Alice , Lisa , Mary
 - exists** returns **true**
- Evaluate **exists** condition on (Alice,111,70K)
 - Subquery returns: Lisa , Mary
 - exists** returns **true**
- Evaluate **exists** condition on (Lisa,222,32K)
 - Subquery returns: empty relation
 - exists** returns **false**
- Evaluate **exists** condition on (Tom,333,56K)
 - Subquery returns: Lisa
 - exists** returns **true**
- Evaluate **exists** condition on (Mary,333,65K)
 - Subquery returns: Lisa
 - exists** returns **true**

Emp x Dept

Emp.ename	Emp.dno	Emp.sal	Dept.dno	Dept.dname	Dept.mgr
Jack	111	81K	111	Sells	Alice
Jack	111	81K	222	Toys	Lisa
Jack	111	81K	333	Electronics	Mary
Alice	111	70K	111	Sells	Alice
Alice	111	70K	222	Toys	Lisa
Alice	111	70K	333	Electronics	Mary
Lisa	222	32K	111	Sells	Alice
Lisa	222	32K	222	Toys	Lisa
Lisa	222	32K	333	Electronics	Mary
Jack	333	81K	111	Sells	Alice
Jack	333	81K	222	Toys	Lisa
Jack	333	81K	333	Electronics	Mary
Tom	333	56K	111	Sells	Alice
Tom	333	56K	222	Toys	Lisa
Tom	333	56K	333	Electronics	Mary
Mary	333	65K	111	Sells	Alice
Mary	333	65K	222	Toys	Lisa
Mary	333	65K	333	Electronics	Mary

E1(ename,dno,salary)

ename	dno	sal
Jack	111	81K
Alice	111	70K
Lisa	222	32K
Tom	333	56K
Mary	333	65K

Result

E1.ename
Jack
Alice
Tom
Mary

Joins

- Joins are expressed implicitly using SELECT-FROM-WHERE clause.
- Alternatively, joins can be expressed using join expressions.
- Different vendors might have different implementations.

Joins

```
SELECT (column_list)
FROM table_name
  [INNER | {LEFT | RIGHT | FULL } OUTER] JOIN table_name
  ON <condition>
WHERE ...
```

- Explicit join semantics needed unless it is an INNER join
- (INNER is default)

Example Schema

```
CREATE TABLE Sailors (  
    sid INTEGER PRIMARY KEY,  
    sname CHAR(20),  
    rating INTEGER,  
    age INT)
```

```
CREATE TABLE Boats (  
    bid INTEGER PRIMARY KEY,  
    bname CHAR (20),  
    color CHAR(10) )
```

```
CREATE TABLE Reserves (  
    sid INTEGER,  
    bid INTEGER,  
    day DATE,  
    PRIMARY KEY (sid, bid, day),  
    FOREIGN KEY (sid) REFERENCES Sailors,  
    FOREIGN KEY (bid) REFERENCES Boats)
```

Inner Join

- Only the rows that match the search conditions are returned.

```
SELECT *  
FROM Sailors INNER JOIN Reserves  
ON Sailors.sid = Reserves.sid
```

- Returns only those sailors who have reserved boats
- How many attributes will the result have?

- Equivalent to:

```
SELECT *  
FROM Sailors, Reserves  
WHERE Sailors.sid = Reserves.sid
```

NATURAL (Inner) Join

- SQL-92 also allows:
 SELECT *
 FROM Sailors NATURAL JOIN Reserves;
- “NATURAL” means equi-join for each pair of attributes with the same name
- Drops the duplicate attributes.
- How many attributes will the result have?

SELECT *

FROM Sailors INNER JOIN Reserves

ON Sailors.sid = Reserves.sid



Sailors(sid,sname,rating,age)

sid	sname	rating	age
22	Dustin	7	45
31	Lubber	8	55
95	Bob	3	60

Reserves(sid,bid,day)

sid	bid	day
22	101	10/10/2012
95	103	11/12/2012

Result

Sailors.sid	sname	rating	age	Reserves.sid	bid	day
22	Dustin	7	45	22	101	10/10/2012
95	Bob	3	60	95	103	11/12/2012

Left Outer Join

- Left Outer Join returns all matched rows, plus all unmatched rows from the table on the left of the join clause
 - use NULL in fields of non-matching tuples

```
SELECT *  
FROM Sailors LEFT OUTER JOIN Reserves  
ON Sailors.sid = Reserves.sid
```

- Returns all sailors & information on whether or not they have reserved boats

SELECT *

FROM Sailors LEFT OUTER JOIN Reserves

ON Sailors.sid = Reserves.sid



Sailors(sid,sname,rating,age)

sid	sname	rating	age
22	Dustin	7	45
31	Lubber	8	55
95	Bob	3	60

Reserves(sid,bid,day)

sid	bid	day
22	101	10/10/2012
95	103	11/12/2012

Result

Sailors.sid	sname	rating	age	Reserves.sid	bid	day
22	Dustin	7	45	22	101	10/10/2012
31	Lubber	8	55	NULL	NULL	NULL
95	Bob	3	60	95	103	11/12/2012

Right Outer Join

- Right Outer Join returns all matched rows, plus all unmatched rows from the table on the right of the join clause

```
SELECT *  
FROM Reserves RIGHT OUTER JOIN Boats  
ON Reserves.bid = Boats.bid
```

- Returns all boats & information on which ones are reserved.

SELECT *

FROM Reserves RIGHT OUTER JOIN Boats

ON Reserves.bid = Boats.bid



Reserves(sid,bid,day)

sid	bid	day
22	101	10/10/2012
95	103	11/12/2012

Boats(bid,bname,color)

bid	bname	color
101	Interlake1	Blue
102	Interlake2	Red
103	Clipper	Green
104	Marine	Red

Result

sid	Reserves.bid	day	Boats.bid	bname	color
22	101	10/10/2012	101	Interlake1	Blue
NULL	NULL	NULL	102	Interlake2	Red
95	103	11/12/2012	103	Clipper	Green
NULL	NULL	NULL	104	Marine	Red

Full Outer Join

- Full Outer Join returns all (matched or unmatched) rows from the tables on both sides of the join clause

```
SELECT *  
FROM Reserves FULL OUTER JOIN Boats  
ON Reserves.bid = Boats.bid
```

SELECT *

FROM Reserves FULL OUTER JOIN Boats

ON Reserves.bid = Boats.bid



Reserves(sid,bid,day)

sid	bid	day
22	101	10/10/2012
95	103	11/12/2012
97	105	1/20/2013

Boats(bid,bname,color)

bid	bname	color
101	Interlake1	Blue
102	Interlake2	Red
103	Clipper	Green
104	Marine	Red

Result

sid	Reserves.bid	day	Boats.bid	bname	color
22	101	10/10/2012	101	Interlake1	Blue
NULL	NULL	NULL	102	Interlake2	Red
95	103	11/12/2012	103	Clipper	Green
97	105	1/20/2013	NULL	NULL	NULL
NULL	NULL	NULL	104	Marine	Red

How about?



```
SELECT *  
FROM Sailors FULL OUTER JOIN Boats  
ON Sailors.sname = Boats.bname
```

Sailors(sid,sname,rating,age)

sid	sname	rating	age
22	Dustin	7	45
31	Lubber	8	55
95	Bob	3	60

Boats(bid,bname,color)

bid	bname	color
101	Interlake1	Blue
102	Interlake2	Red
103	Clipper	Green
104	Marine	Red

How about?



SELECT *

FROM Sailors

RIGHT OUTER JOIN Reserves ON Sailors.sid = Reserves.sid

LEFT OUTER JOIN Boats ON Reserves.bid = Boats.bid

Sailors(sid,sname,rating,age)

sid	sname	rating	age
22	Dustin	7	45
31	Lubber	8	55
95	Bob	3	60

Reserves(sid,bid,day)

sid	bid	day
22	101	10/10/2012
95	103	11/12/2012
97	105	1/20/2013

Boats(bid,bname,color)

bid	bname	color
101	Interlake1	Blue
102	Interlake2	Red
103	Clipper	Green
104	Marine	Red

Join Summary

- R INNER JOIN S ON <condition>
- R FULL OUTER JOIN S ON <condition>
- R LEFT OUTER JOIN S ON <condition>
- R RIGHT OUTER JOIN S ON <condition>
- R NATURAL JOIN S
- R NATURAL FULL OUTER JOIN S
- R NATURAL LEFT OUTER JOIN S
- R NATURAL RIGHT OUTER JOIN S

Again: Different vendors might have different implementations.

More on Subqueries

- **Subqueries in the FROM statement**

- **Example:** “List the department dno, total salary, and number of employees for all departments with more than 1 employee.”

```
SELECT dno, SUM(sal), COUNT(ename)
FROM Emp
GROUP BY dno
HAVING COUNT(ename)>1;
```

```
SELECT Temp.dno, Temp.totalSal, Temp.numEmp
FROM (
    SELECT dno, SUM(sal) as totalSal , COUNT(ename) as numEmp
    FROM Emp
    GROUP BY dno
) AS Temp
WHERE Temp.numEmp > 1;
```

Both queries return the same result.

More on Subqueries

- **Subqueries to aggregate in multiple stages**
 - **Example:** “Find the department dno that has the max total salary amount among all departments.”

```
SELECT Temp.dno, MAX(Temp.totalSal)
FROM (
    SELECT dno, SUM(sal) as totalSal
    FROM Emp
    GROUP BY dno
) AS Temp;
```

More on Subqueries

- **Joining Subqueries**

- **Example:** “Find the employees in the departments whose total salary amount is more than \$200K”

```
SELECT Emp.ename
FROM (
    SELECT dno, SUM(sal) as totalSal
    FROM Emp
    GROUP BY dno
) AS Temp INNER JOIN Emp ON Temp.dno=Emp.dno
WHERE Temp.totalSal>200000;
```

```
SELECT Emp.ename
FROM Emp
WHERE Emp.dno IN (
    SELECT dno
    FROM Emp
    GROUP BY dno
    HAVING SUM(sal) >200000);
```

More Subquery Examples

Emp(ename, dno, salary)

Dept(dno, dname, mgr)

Emp(dno) is a FK referencing Dept(dno)

Dept(mgr) is a FK referencing Emp(ename)

Query 1: *“Find the employees whose salary is greater than overall average salary”.*

Query 2: *“Find the employees whose avg salary is greater than the average salary in the department they work in”.*

Query 3: *“Find the department dno which has the minimum average salary. ”*

More Subquery Examples

Suppliers(sid, sname, address)

Parts(pid, pname, color)

Catalog(sid, pid, cost)

Catalog(sid) is a FK referencing Suppliers(sid)

Catalog(pid) is a FK referencing Parts(pid)

Query 1: *“Find the distinct “sids” of suppliers who charge more for some part than the average cost of that part (averaged over all the suppliers who supply that part)”.*

Query 2: *“Find the distinct “sids” of suppliers who supply only red parts”.*